

Learning Implied Volatility, Not Prices: A Neural-Operator Encoder with a Differentiable Black–Scholes Decoder for SPX Option Pricing

Zefang Li
zli299@jhu.edu

Abstract

The Black–Scholes equation (BS) gives accurate prices for European options when given the implied volatility. However, how much per-contract conditioning helps remains unquantified. Here, we adopt neural operators to learn a per-contract implied volatility $\hat{\sigma}$ and feed it through the closed-form BS equation, so a plain data loss on the log normalized price $\log(\frac{V}{K})$ backpropagates through the analytical pricer, and the Greeks come from the closed-form formula rather than a learned approximation. A 2D Fourier Neural Operator encodes the market state from one of three inputs (implied-volatility surface, SPX price history, or VIX history), and a per-query head emits $\hat{\sigma}$ from each contract’s log-moneyness and maturity. On quote-quality-filtered SPX options (OptionMetrics via WRDS) split 80/10/10 by trade date, three results follow. Per-query conditioning matters most: one scalar $\hat{\sigma}$ per market state rather than per contract collapses put $R^2(\log)$ from 0.978 to about -0.32 and $R^2(\text{price})$ to -3.090 . The surface-conditioned configuration prices best, $R^2(\log) \approx 0.98\text{--}0.99$ against $\approx 0.90\text{--}0.94$ for the historical ones. Autograd partial greeks through the decoder match closed-form BS at the predicted $\hat{\sigma}$ to a maximum absolute discrepancy of 2×10^{-14} (delta) and 2×10^{-15} (vega). Two qualifications apply: the surface experiment is same-day cross-sectional reconstruction, not forecasting, and the comparison is of whole input–encoder configurations, not isolated inputs.

1 Introduction

Pricing an option is, in principle, a map from market state and contract terms to a price. The common neural approach regresses that price (or an implied volatility) directly with a general function approximator, a line of work that begins with Hutchinson et al. [13] and is surveyed by Ruf and Wang [20]. Regressing the price directly throws away structure that is already known in closed form: under the standard Black–Scholes assumptions, the price is an explicit, differentiable function of one hard-to-observe quantity, the implied volatility σ , conditional on moneyness, maturity, and the dataset-supplied rate and dividend-yield inputs. However, that literature leaves two empirical questions open about such an architecture: how much a per-contract $\hat{\sigma}$ buys relative to a market-state-only scalar $\hat{\sigma}$, and which of three market-state inputs (an implied-volatility surface, an SPX price history, or a VIX history) prices best when the architecture, dataset, and date split are held fixed. Neither has been quantified side by side on SPX options [20]. Here, we train a network to predict only σ —more precisely, a per-contract estimate $\hat{\sigma}$ conditioned on the broader market state—and recover the price analytically through the BS formula. Because BS is composed of differentiable primitives ($\exp$, erf , $\sqrt{\cdot}$), the training loss on $\log(\frac{V}{K})$ back-propagates through the closed-form decoder into the network exactly as it would through any other differentiable layer; no learned pricing surrogate is needed.

Two structural properties follow from this design, and we state them at the precision the checks in this paper actually support. First, conditional on a fixed predicted $\hat{\sigma}$, the decoder returns the

closed-form BS price, respects its single-contract static no-arbitrage bounds, and has the corresponding analytical partial derivatives. This is a pointwise, frozen- $\hat{\sigma}$ Black–Scholes-consistency property, not a claim about the total learned map: $\hat{\sigma}$ itself varies with the query $(\log(\frac{S}{K}), T)$, so the composed map from market state and contract terms to price need not satisfy the BS PDE, whose derivation assumes a volatility that does not depend on strike or maturity [8]. Nor is it a claim that the full predicted surface is arbitrage-free—because $\hat{\sigma}$ is fit independently for each query, nothing in the architecture guarantees calendar-spread or butterfly consistency across the predicted (K, T) surface, and we do not check this here. Second, every greek we report is the closed-form BS partial derivative evaluated at the predicted $\hat{\sigma}$, holding $\hat{\sigma}$ fixed—an analytic partial greek, not the total sensitivity of the learned system, which would include an additional $\partial\hat{\sigma}/\partial S \cdot \text{vega}$ term that a frozen- $\hat{\sigma}$ check cannot see.

Contributions. This is an empirical, application-focused study rather than a new-method paper: the differentiable-BS-decoder idea is not novel to the machine learning or quantitative finance communities (Section 2). We make three specific, evaluated claims.

1. Per-query volatility conditioning is what makes the smile/skew learnable. We compare the full per-query architecture against a scalar- $\hat{\sigma}$ ablation that predicts one volatility per market state (flat-vol BS, no smile or skew) with everything else held fixed. The gap is large and, for puts, qualitative: removing per-query conditioning drives put $R^2(\text{price})$ negative (Section 6.3).
2. We quantify which of three candidate market-state input–encoder configurations prices best. Comparing an implied-volatility-surface encoder against SPX-price-history and VIX-history encoders, the surface configuration wins by a wide and consistent margin on both option types and both head architectures (Section 6.1). Two qualifications travel with this result. (a) The surface branch is fed OptionMetrics’ own same-day standardized volatility surface, derived from that trading day’s option quotes; the date-based split rules out cross-date leakage but not this same-day dependence, so that branch demonstrates same-day cross-sectional reconstruction conditional on an observed surface rather than independent price discovery. The SPX-history and VIX-history tensors are 21-trading-day windows ending on the target trade date and therefore include that day’s underlying and VIX closes; their same-day dependence on the option market is weaker and far lower-dimensional than the surface branch’s (SPX history has none directly, only the same-day underlying close; VIX is a single scalar index that Cboe computes from contemporaneous SPX option quotations), but it is not zero. (b) The three encoders are not parameter-matched—the flattened FNO output is projected to the latent vector by two dense layers, and the first of them, whose input width scales with the grid, dominates the count: it holds 766,080 parameters (weights and bias) for the 11×17 surface grid, 430,208 for the 21×5 SPX-history grid, and 344,192 for the 21×4 VIX-history grid, so input content, timing, normalization, grid size, and capacity all vary together and the comparison is between whole configurations, not raw inputs (Section 7).
3. The BS decoder gives verified derivative consistency at the predicted $\hat{\sigma}$. We compare autograd-computed delta and vega through the decoder against the closed-form BS expressions at the same $\hat{\sigma}$ and find agreement at float64 machine precision, confirming the decoder is implemented and differentiated correctly (Section 6.6).

We additionally report a comparison between the two per-query head designs used throughout (an MLP head and a DeepONet branch–trunk readout). The observed differences are small relative to the seed-to-seed variation we measure, and three seeds are insufficient to rank the two architectures, so we do not treat this as a fourth contribution or a ranking claim (Section 6.2).

2 Related work

Option pricing theory. The closed-form Black–Scholes formula [2] and its rational-pricing derivation [16] are the analytical decoder we build on; we make no claim to extend or revise the theory itself, only to use it as a differentiable output layer.

Neural networks for option pricing. Hutchinson et al. [13] introduced nonparametric learning networks for pricing and hedging derivative securities, and Ruf and Wang [20] survey the three decades of work that followed. Culkin and Das [6] trained a feed-forward network to reproduce Black–Scholes prices directly, establishing that a network can learn the BS pricing function to high accuracy as a numerical approximation. Horvath et al. [12] instead used networks to accelerate calibration of (rough) stochastic volatility models against market implied volatility surfaces. Our design differs from all of these in keeping the BS formula itself analytical and training the network only to output its one free parameter, $\hat{\sigma}$, per contract; we do not train a network to approximate BS itself, and we do not calibrate a parametric stochastic volatility model.

Implied-volatility-surface-conditioned neural pricing. Closest in spirit to our surface branch, Gonon et al. [11] use a neural operator to map raw, irregularly observed option quotes directly to a smoothed implied-volatility surface, and Ding et al. [7] compress observed market volatility surfaces with a variational autoencoder and price contracts from the resulting latent code together with strike and maturity. Ackerer et al. [1] fit the implied-volatility surface itself with a neural network trained on the observed quotes, adding soft no-arbitrage penalties on the predicted surface to the loss. All three condition on contemporaneous option-market information, as our surface-branch experiment does. Our contribution relative to them is not the conditioning itself but the analytical BS output layer (the network emits only $\hat{\sigma}$) and the side-by-side comparison of that surface input against purely historical inputs on the same architecture, dataset, and split.

Physics-informed and structure-constrained approaches. Physics-informed neural networks [19] enforce a governing PDE via a soft penalty term added to the training loss, originally in a physics setting and subsequently adapted to option pricing and calibration problems by other authors. Our architecture reaches a related but weaker property—at a frozen predicted $\hat{\sigma}$, the closed-form BS price and its single-contract static no-arbitrage bounds at each queried contract—structurally, by construction of the decoder, rather than via an auxiliary loss term; Section 3.4 states precisely what this guarantees and does not guarantee.

Neural operators. The market-state encoder is a 2-D Fourier Neural Operator [14]. We also use a DeepONet-style branch–trunk readout [15] as an alternative per-query head, dotting a market-state branch vector against a trunk network evaluated at each contract’s (log-moneyness, T) query. Noguer i Alonso [17] train operator-learning surrogates for finance against both prices and directional derivatives, and draw the same distinction we rely on in Sections 3.4 and 6.6: a surrogate that matches prices need not match their sensitivities, so derivative consistency is a separate property that must be established on its own terms rather than inferred from pricing accuracy.

Classical implied-volatility surface modeling. SVI [9] and its arbitrage-free extension SSVI [10] are standard practitioner parameterizations of the implied-volatility smile; we use spline and nearest-neighbor interpolation of the observed surface as one of our practitioner baselines (Section 5.2) rather than fitting SVI directly.

Quote-quality filtering. Excluding option observations that violate static no-arbitrage bounds prior to model fitting has precedent in the literature: Büchner and Kelly [3] filter option return panels on data-quality grounds in an asset-pricing context, and Cohen et al. [5] both detect and repair arbitrage violations in traded option price data, cautioning that outright removal (as opposed to repair) discards information. Our dataset filter (Section 4.3) follows the removal approach, with the caveat these authors raise explicitly noted.

3 Method

3.1 The inverse implied-volatility problem

Given an observed normalized option price $\frac{V}{K}$ for a contract with log-moneyness $m = \log(\frac{S}{K})$, maturity T , risk-free rate r , and dividend yield q , the implied volatility σ is the value that makes the BS formula reproduce the observed price. We train an estimator $\hat{\sigma}$ of this quantity, conditioned on the broader market state, and price through BS. The training target is $\log(\frac{V}{K})$ rather than $\frac{V}{K}$ itself, since option prices span several orders of magnitude across moneyness and maturity.

3.2 Market-state encoder

The market state is encoded by a 2-D Fourier Neural Operator [14]: four spectral convolution blocks with 1×1 -convolution skip connections and SiLU activations, flattened and projected to a 128-dimensional latent vector. We evaluate three candidate market-state inputs, one per encoder instance, each reshaped from its native 1-D form to a 2-D grid inside the encoder.

The implied-volatility surface input is OptionMetrics’ standardized volatility surface [18] for the same trade date, on an 11×17 (maturity $\times$ delta) grid. This input is derived from that day’s own option quotes (Section 4.1), so this branch is conditioned on contemporaneous cross-sectional information rather than on strictly prior information. The SPX price history input is 21 trading days of SPX OHLCV data, a 21×5 grid, each row divided by the window’s final close (and volume by the final volume); the window ends on the target trade date rather than the day before it, so the final row is that day’s own OHLCV and that day’s close is the normalization denominator for the whole tensor. The VIX history input is 21 trading days of CBOE VIX OHLC data, a 21×4 grid, likewise ending on the target trade date and normalized the same way: each value is divided by the window’s final (i.e. same-day) VIX close. This representation is therefore scale-free and retains only the shape of the recent VIX path; the absolute VIX level is not visible to that encoder.

Neither historical input is a strictly-prior-information input: both windows include the target trade date, and VIX is itself computed by Cboe from a strip of contemporaneous SPX option quotations [4], so even the VIX branch is not free of same-day option-market information. What separates these branches from the surface branch is the amount and directness of that information—a single scalar index (further reduced to a shape by the normalization) versus the full standardized 11×17 implied volatility surface—not its presence or absence.

The flattened FNO output is projected to the latent vector by two dense layers, a flatten-to-128 layer whose input width is proportional to the grid size followed by a fixed 128-to-latent layer, so the three encoders differ in parameter count. The first layer dominates that difference and is the one the published counts refer to: including its bias it holds 766,080 parameters for the surface grid, 430,208 for SPX history, and 344,192 for VIX history, while the second layer is identical across all three branches. The branch comparison in Section 6.1 is consequently not a capacity-controlled comparison, and we read it accordingly (Section 7).

3.3 Per-query volatility head

The market-state latent vector feeds a per-query head that additionally takes each contract’s (log-moneyness, T) query and outputs $\hat{\sigma}$ for that specific contract, not one scalar per market state. We compare two interchangeable head designs, both with hidden width 128 acting on the same 128-dimensional latent vector. Architecture A (MLP head) is an MLP on the concatenation of the latent vector and the query (log-moneyness, T), followed by a softplus to keep $\hat{\sigma}$ positive. Architecture B (DeepONet) is a DeepONet-style branch–trunk readout [15]: the FNO output is the branch vector, a trunk MLP embeds (log-moneyness, T), and $\hat{\sigma}$ is a softplus of their dot product plus a bias, with a small positive floor added for numerical stability.

Per-query conditioning is the central design choice of this paper: because $\hat{\sigma}$ depends on the specific contract being priced, not only on the market state, the model can in principle recover the volatility smile/skew, which a market-state-only $\hat{\sigma}$ cannot. Section 6.3 isolates the effect of this choice.

3.4 Differentiable Black–Scholes decoder

A closed-form BS pricing function maps (log-moneyness, $T, r, q, \hat{\sigma}$) to the normalized price $\frac{V}{K}$ using only differentiable primitives ($\exp$, erf , $\sqrt{\cdot}$). Gradients of the data loss on $\log(\frac{V}{K})$ flow through this decoder into the encoder and head exactly as through any other differentiable layer; there is no learned or neural pricing surrogate anywhere in the forward pass. Figure 1 summarizes the full forward pass.

Two properties follow directly from using an analytical decoder. First, every reported greek is the closed-form BS partial derivative evaluated at the predicted $\hat{\sigma}$, holding $\hat{\sigma}$ fixed. We verify this numerically in Section 6.6 by comparing autograd through the decoder against the closed-form expressions; this checks decoder correctness, not the total (sticky-strike) sensitivity of the learned system, which would additionally involve $\partial\hat{\sigma}/\partial S$. Second, conditional on a fixed predicted $\hat{\sigma}$, the decoder returns the closed-form BS price, respects its single-contract static no-arbitrage bounds, and has the corresponding analytical partial derivatives—a property of the frozen- $\hat{\sigma}$ mapping only. It is not a claim that the total learned map solves the BS PDE: the predicted $\hat{\sigma}$ is a function of the query ($\log(\frac{S}{K}), T$), whereas the BS PDE is derived under a volatility that does not vary with strike or maturity, so the composed map need not satisfy it. Nor does the pointwise property guarantee calendar-spread or butterfly consistency across the full predicted (K, T) surface, since $\hat{\sigma}$ is estimated independently per query; we do not run a surface-wide arbitrage violation-rate check in this paper.

A third property is a limit rather than a guarantee. As $T \rightarrow 0$ the Black–Scholes price converges to the discounted forward intrinsic value for every σ , so at exact expiry the theoretical price carries no information about volatility and the implied volatility is unidentified there. The settlement-aware maturity filter of Section 4.2, which retains only contracts with strictly more than 24 hours remaining to settlement, excludes this case from the sample. The implementation is not exactly $\hat{\sigma}$ -independent at $T=0$: it evaluates $\sqrt{\text{clamp}(T, 10^{-10})}$, clamping the maturity before taking the square root rather than clamping the square root itself, and therefore retains a residual $\hat{\sigma}$ dependence of order 10^{-5} in a vanishing band around the money: a synthetic exact- $T=0$ probe of the decoder finds a maximum price spread over $\hat{\sigma}$ of 1.2×10^{-5} in $\frac{V}{K}$ units (results/diagnostics_tail_v5.json). No contract in this dataset reaches that regime.




Figure 1: Forward pass. One of three candidate market-state inputs (implied-volatility surface, SPX price history, or VIX history) is encoded by a 2-D FNO into a shared latent vector. A per-query volatility head (architecture A: MLP head, or architecture B: DeepONet branch-trunk) combines the latent vector with a contract-specific query ($\log(\frac{S}{K}), T$) to produce $\hat{\sigma}$. The analytical Black-Scholes decoder takes five inputs— $\log(\frac{S}{K}), T$, the rate r , the dividend yield q , and $\hat{\sigma}$ —and returns the normalized price $\frac{V}{K}$; the first four are observed contract and market data routed to the decoder directly, and only $\hat{\sigma}$ comes from the network. The pricing function itself is fixed and differentiable.

3.5 Loss

Training minimizes plain mean-squared error on $\log(\frac{V}{K})$, computed as $\log(\text{clamp}(\frac{V}{K}, 10^{-8}))$ for numerical stability. There is no auxiliary PDE, arbitrage, or BS-anchor penalty term. We do not claim such terms would be redundant: the pointwise consistency of Section 3.4 says nothing about calendar-spread or butterfly consistency across the predicted surface, which is exactly what a soft arbitrage penalty of the kind used by Ackerer et al. [1] would target. The plain data loss is used because it was sufficient for the questions this paper asks and keeps the training objective free of tuned penalty weights; adding and evaluating a surface-wide calendar-spread and butterfly penalty on the predicted (K, T) grid is left to future work.

4 Data

We use SPX index options from OptionMetrics, accessed via WRDS, covering 2015-02-02 through 2025-08-29. European exercise removes the early-exercise complication, and BS is used as the analytical decoder. Contracts are filtered to those with nonzero traded volume.

4.1 The implied-volatility-surface input and its same-day dependence

The surface branch input is OptionMetrics’ standardized volatility surface [18]: an 11×17 grid of interpolated implied volatilities over eleven fixed maturities (10–730 calendar days) and seventeen fixed BS spot-delta points (+10 to +90 for calls, −90 to −10 for puts). Preprocessing joins this surface to the option targets on the same trade date. Because OptionMetrics [18] construct the surface by interpolating that day’s own option implied volatilities, the input can share information with the population of quotes being priced: the date-based split described below assigns each trade date entirely to one split, but it does not and cannot remove this within-date dependence. Accordingly, the surface-branch experiment should be read as out-of-time (date-split), same-day cross-sectional reconstruction conditional on an observed surface—a nowcasting task—and not as independent price discovery or as forecasting.

The SPX-history and VIX-history inputs are built from a 21-trading-day OHLC(V) window that ends on the target trade date, so they are not lagged-only inputs: the final row is the trade date’s own bar, and that day’s close is the denominator used to normalize the entire tensor. Their same-day dependence on the option market is nevertheless much weaker than the surface branch’s. SPX history sees only the same-day underlying price; VIX history sees a single scalar index, which Cboe computes from a weighted strip of contemporaneous SPX option bid/ask quotations [4], further reduced by this representation to a path shape. The ordering in directness of same-day option-market information is therefore surface > VIX history > SPX history, with only the surface branch receiving the day’s full cross-section of implied volatilities.

4.2 Maturity and contract-identity handling

Each contract’s maturity is computed on a settlement-aware basis: AM-settled contracts lose the 6.5 hours between the prior end-of-day quote and the 9:30 ET opening settlement print, and we retain only contracts with strictly more than 24 hours remaining to settlement. Contract identity (including AM/PM settlement, symbol, and OptionMetrics option ID) is preserved through preprocessing, so distinct contracts are never collapsed onto identical model inputs.

4.3 Quote-quality filter

We apply a deterministic, zero-tolerance static no-arbitrage filter to the observed midpoint of every contract. In normalized $\frac{V}{K}$ units, with forward price $\text{fwd} = Me^{-qT}$ and discount factor $\text{disc} = e^{-rT}$, a call midpoint must lie in $[\max(\text{fwd} - \text{disc}, 0), \text{fwd}]$ and a put midpoint in $[\max(\text{disc} - \text{fwd}, 0), \text{disc}]$, computed in float64 before conversion to the stored float32 representation. Any row whose midpoint falls outside this bound under the dataset’s own supplied rate, dividend-yield, and settlement inputs is removed. Table 1 reports the resulting row counts: this filter removes 2.21% of call rows and 0.82% of put rows overall, with a small train-versus-test differential (well under one percentage point on both option types). We preferred this rule to spread-based filtering alternatives, which would have shifted train and test distributions unevenly given the secular tightening of SPX quoted spreads from 2015 to 2025. We do not claim that every removed midpoint reflects a realizable market arbitrage: the rate, dividend-yield, and settlement timestamp inputs used to compute the bound are themselves modeling choices in this dataset, not observed riskless-rate trades, so the filter’s claim is that the BS decoder cannot reach these midpoints at any $\hat{\sigma}$ given those inputs.

After filtering, the dataset validation pass finds zero static-bound violations, no non-finite values, and a minimum maturity of 1.73 days across all 4,846,268 call and 7,925,349 put rows, and a separate duplicate-input check finds 0.00% duplicate-input rows within every split, over the full dataset, and across split boundaries, so every contract maps to a distinct model input. No contract is quoted exactly on its own expiry date under the settlement-aware maturity cut, so the $T=0$ degeneracy discussed in Section 3.4, where the theoretical Black–Scholes price collapses to intrinsic value and the implied volatility is unidentified, does not occur in this sample.

4.4 Train/validation/test split

The 80/10/10 split is by unique trade date, not by row: all contracts quoted on a given date are assigned entirely to one split. This is a deliberate, time-ordered split with disjoint target dates: no contract’s quote date appears in more than one split. It is not a claim that no information whatsoever crosses a split boundary. The 21-trading-day trailing SPX and VIX feature windows reach back across those boundaries—a test-split contract’s history window includes dates assigned to validation or train—which is a normal feature of out-of-time evaluation with trailing-window

Table 1: Row counts by option type and split, before and after the quote-quality filter of Section 4.3. Split date boundaries are identical before and after filtering (the filter removes individual rows, never a trade date): train 2,128 dates (2015-02-02 – 2023-07-17), validation 266 dates (2023-07-18 – 2024-08-06), test 267 dates (2024-08-07 – 2025-08-29).

Option	Split	Unfiltered rows	Filtered rows	Removed	Removed %
Call	Train	3,360,823	3,297,722	63,101	1.878%
Call	Validation	754,276	729,204	25,072	3.324%
Call	Test	840,493	819,342	21,151	2.516%
Call	Total	4,955,592	4,846,268	109,324	2.206%
Put	Train	5,370,629	5,323,327	47,302	0.881%
Put	Validation	1,229,955	1,221,643	8,312	0.676%
Put	Test	1,390,470	1,380,379	10,091	0.726%
Put	Total	7,991,054	7,925,349	65,705	0.822%

inputs, not a defect of this split. Nor does the split address dependence within a date: the surface branch input is a same-day construct (Section 4.1), so the split makes the evaluation out-of-time but leaves that branch’s task a same-day cross-sectional one.

4.5 What is not reconstructed

The raw OptionMetrics CSV row count and a per-rule breakdown of upstream attrition (secid/index filters, date-range restriction, VIX/spot join losses) are not recorded as a committed artifact in this pipeline; only the post-preprocessing parquet total (13,595,211 rows, both option types combined, prior to the maturity cut) is. We report the retention that is directly measured (Table 1) and do not claim a finer breakdown than the pipeline actually recorded.

4.6 Data availability

The SPX options data used in this study are licensed OptionMetrics data accessed through a WRDS subscription; CBOE VIX history is publicly available from CBOE. WRDS/OptionMetrics subscription data is, under standard WRDS terms, not freely redistributable, so we release the preprocessing and model code but not the underlying dataset; a WRDS subscriber with OptionMetrics access can reproduce the dataset used here from the raw SPX options tables (secid 108105) and CBOE VIX history using the released preprocessing code. We state the general norm for WRDS-licensed data here; exact redistribution terms should be confirmed against the specific WRDS/OptionMetrics license in force at submission time.

5 Experiments

5.1 Core training runs

We train all combinations of $\{\text{call, put}\} \times \{\text{IV surface, SPX history, VIX history}\} \times \{\text{architecture A, architecture B}\}$, twelve runs in total, on the filtered dataset of Section 4 with a shared training configuration: a first phase of Adam optimization (100 epochs, learning rate 10^{-3} , reduce-on-plateau scheduling) followed by a fine-tuning phase (50 epochs, learning rate 10^{-5}); batch size 256; gradient-norm clipping at 1.0; early stopping on validation MSE with patience 5 in

each phase. All twelve runs use seed 42; the four headline IV-surface runs (call/put $\times$ architecture A/B) are additionally trained at two more seeds (43, 44) to give a dispersion estimate for the architecture comparison and for the surface branch’s own stability in Section 6. That replication does not cover the branch-vs-branch comparison: the SPX-history and VIX-history runs remain single-seed, and we do not report or imply a variance estimate for them. Two of those eight single-seed history runs, call SPX history with architecture A and put SPX history with architecture B, reached their best validation loss at epoch 1 of the 16 epochs recorded and never improved on it (`train_model_v3/call/results_spot_history_v5/loss_history.txt`, `train_model_v3/put/results_spot_history_don_v5/loss_history.txt`), so the checkpoints carried into those two cells of Table 2 are effectively one-epoch models. The single-seed history-branch results should therefore be read as potentially optimization-sensitive rather than as tuned ceilings for those configurations. Test-set metrics reported throughout are computed on the normalized price $\frac{V}{K}$, not raw dollar prices, and are averaged over the full test split (no subsampling) by a held-out evaluation script loaded from each run’s best checkpoint by validation loss.

5.2 Baselines

We compare the trained models against three baselines with no learned per-query conditioning: the zero-volatility limit (discounted forward intrinsic), which prices a call at $\max(\text{fwd} - \text{disc}, 0)$ and a put at $\max(\text{disc} - \text{fwd}, 0)$ in $\frac{V}{K}$ units and fits no parameter at all; a single constant $\hat{\sigma}$ per option type (flat-vol BS), fitted under two separate objectives, minimum MSE on $\log(\frac{V}{K})$ and minimum MSE on $\frac{V}{K}$, because the two objectives select materially different constants; and classical implied-volatility-surface interpolation (nearest-neighbor, bilinear, and cubic-spline fits to the observed 11×17 surface, then BS). All three are evaluated on the same filtered test split as the models. Every fitted baseline constant reported in this paper, the two constant- $\hat{\sigma}$ values and every scale factor alike, is fitted on one fixed uniform 300,000-row subsample of the training split drawn with seed 0 (`results/baselines_v5.json`), not on the full training split. The same sweep also evaluates a family of VIX-level baselines that set $\hat{\sigma}$ to the day’s VIX level times a single scale factor fitted on that same train subsample; we do not tabulate these as headline comparators, since they use a market observable that is not one of the three branch inputs, but we quote the price-space-fitted variant, `vix_level_scaled_pxfit`, in Section 6.1, where it bears on how the VIX result should be read.

Surface-interpolation details. The surface baseline queries the standardized grid on a regular (calendar-days, delta-percent) lattice. Two points deserve explicit acknowledgment. First, the surface’s moneyness axis is BS spot delta, which is itself a Black–Scholes-derived quantity computed from a volatility, so a contract’s grid coordinate depends on the volatility being looked up. We resolve this with a damped fixed-point iteration ($\sigma \leftarrow \text{surface}(T, \delta(\sigma))$, damping 0.7, initialized at the at-the-money $|\delta| = 50$ column, tolerance 10^{-6}), and report the non-convergence count as a diagnostic. This baseline is consequently not model-free in the strict sense (it inherits the BS delta convention), which is appropriate here since it is compared against a pricer that also decodes through BS. Second, queries outside the observed grid (tenors below 10 or above 730 days, deltas outside the 10–90 band) are clamped to the nearest grid boundary rather than extrapolated, so the baseline never invents surface values in regions OptionMetrics does not report. On the test split this clamping reaches a large minority of rows: 227,607 call and 602,724 put rows are clamped on the delta axis, and 170,616 call and 274,497 put rows on the tenor axis. We report all three interpolation orders—nearest, bilinear, and cubic spline—rather than choosing one; cubic spline is the one carried into Table 5 and Table 6 because it is the strongest of the three on both option

types (calls, $R^2(\log)$: 0.7628 nearest, 0.7641 bilinear, 0.7642 spline), and because it is the smoothest, which is the fairest comparator against a smooth learned surface. The three differ only marginally, so no conclusion in this paper turns on that choice.

5.3 Ablation: per-query vs. scalar $\hat{\sigma}$

To isolate the value of per-query conditioning specifically, as opposed to any other difference between architectures or inputs, we train a restricted variant of the IV-surface, architecture-A model in which $\hat{\sigma}$ is produced from the market-state latent alone, with no dependence on the contract query. Everything else (encoder, decoder, training schedule, seed) is held fixed. This variant is not a candidate model; it exists only to quantify the effect of removing per-query conditioning. Like the ablation’s per-query reference run, it is trained at a single seed (42) and is not replicated, so the reported magnitudes describe that one run rather than an averaged estimate. The representational restriction itself— $\hat{\sigma}$ constant across all contracts sharing a market state, and therefore unable to express a smile or skew—is seed-independent; that much follows from the architecture alone. The direction and magnitude of the observed test metrics, however, come from a single unreplicated seed and are not guaranteed by the restriction itself.

5.4 Greeks consistency check

For the trained IV-surface, architecture-A models, we sample 2,000 test contracts per option type and compute delta ($\frac{\partial \text{price}}{\partial S}$) and vega in two ways: by autograd through the Black–Scholes decoder only, and by the closed-form BS expressions, both evaluated at the same predicted $\hat{\sigma}$. The model’s $\hat{\sigma}$ is explicitly detached before the autograd pass, so no gradient flows back through the head or the FNO encoder; the quantity differentiated is the closed-form BS price as a function of (S, T) at a frozen $\hat{\sigma}$, not the encoder–head–decoder composition. Agreement between the two checks that the decoder is implemented and differentiated correctly, not that the learned $\hat{\sigma}$ itself is independently validated.

6 Results

6.1 Main results across market-state inputs and architectures

Table 2 reports test-set metrics for all twelve core runs. Because the settlement-aware maturity filter removes every contract with $T \leq 1$ day, $R^2(\log, \text{full})$ and $R^2(\log, T > 1 \text{ day})$ coincide numerically on this dataset, so we report a single $\log\text{-}R^2$ column.

The ranking across the three input–encoder configurations is consistent across both option types and both architectures: IV surface $\gg$ SPX history $\approx$ VIX history on the headline log metric. Calls score 0.990–0.991 (surface) against 0.922–0.944 (SPX history) and 0.921–0.940 (VIX history); puts score 0.978–0.981 (surface) against 0.902–0.909 (SPX history) and 0.918–0.921 (VIX history). SPX-history and VIX-history inputs are close to each other on both option types.

Three qualifications bound what this ranking supports. First, the surface branch is conditioned on the same-day standardized surface (Section 4.1), so its advantage is measured in a same-day cross-sectional reconstruction setting; the historical branches, whose 21-day windows end on the trade date but carry only that day’s underlying close (SPX history) or a single same-day VIX level reduced to a path shape (VIX history), are solving a substantially harder problem, so the gap is not a like-for-like measure of how much “signal” the two input families contain. Second, the encoders are not parameter-matched (766,080 / 430,208 / 344,192 parameters in the first,

Table 2: Test-set metrics for all twelve core runs (seed 42). $R^2(\log)$ is on $\log(\frac{V}{K})$; $R^2(\text{price})$ is on $\frac{V}{K}$.

Option	Branch	Arch.	$R^2(\text{price})$	$R^2(\log)$	RMSE(log)	MAE(price)
Call	IV surface	A	0.999977	0.991044	0.221894	0.000641
Call	IV surface	B	0.999914	0.990290	0.231041	0.000739
Call	SPX history	A	0.999812	0.922020	0.654761	0.003569
Call	SPX history	B	0.999181	0.944390	0.552923	0.002769
Call	VIX history	A	0.999848	0.921333	0.657639	0.003114
Call	VIX history	B	0.998125	0.939874	0.574940	0.002985
Put	IV surface	A	0.992553	0.977898	0.294136	0.000823
Put	IV surface	B	0.993780	0.981298	0.270566	0.000659
Put	SPX history	A	0.952078	0.909410	0.595489	0.002484
Put	SPX history	B	0.942330	0.901553	0.620777	0.002847
Put	VIX history	A	0.962518	0.920722	0.557070	0.002429
Put	VIX history	B	0.961502	0.917777	0.567323	0.002477

Table 3: Three-seed replication (seeds 42, 43, 44), IV-surface branch only.

Option	Arch.	$R^2(\text{price})$ mean $\pm$ SD	$R^2(\log)$ mean $\pm$ SD	RMSE(log) mean $\pm$ SD
Call	A	0.999972 ± 0.000005	0.990996 ± 0.000073	0.222493 ± 0.000896
Call	B	0.999865 ± 0.000131	0.991075 ± 0.000716	0.221388 ± 0.008838
Put	A	0.994660 ± 0.001898	0.979099 ± 0.001424	0.285924 ± 0.009813
Put	B	0.995392 ± 0.001403	0.981738 ± 0.000611	0.267345 ± 0.004491

grid-dependent layer of the flattening projection), and they also differ in grid size, spectral modes, and input normalization, so input content, timing, and capacity vary together and what is ranked is three complete configurations rather than three raw inputs. Third, the VIX-history tensor is divided by the window’s final VIX close, which removes the absolute level of VIX and leaves only the shape of the 21-day path; the correct reading of the VIX result is that this close-normalized representation underperforms the surface, not that VIX history is uninformative in general. The `vix_level_scaled_pxfit` baseline, which uses the raw VIX level scaled by a single train-subsample-fitted constant, attains the highest call-side $R^2(\log)$ of any non-neural baseline we evaluated, 0.7904 (price-space $R^2 = 0.999809$), above the 0.7642 of spline interpolation of the observed surface in Table 5.

Within these qualifications, the ordering itself is stable: it holds on both option types and both head architectures. Three-seed replication (Table 3) exists only for the IV-surface cells, and there the seed-42 gap to the historical configurations is much larger than the surface branch’s own dispersion. Call architecture A’s across-seed standard deviation is 0.000073 against a gap of $0.991 - 0.922 \approx 0.07$ over SPX-history architecture A, and the least favorable of the four replicated cells still leaves the gap roughly 40 times the corresponding standard deviation. This bounds the stability of the surface branch’s own score only; the SPX-history and VIX-history configurations are single-seed, so the branch-versus-branch comparison itself has not been replicated.

6.2 Architecture comparison: MLP head (A) vs. DeepONet (B)

Table 3 reports mean $\pm$ standard deviation across seeds 42, 43, and 44 for the four IV-surface cells, the only cells with multi-seed replication.

Table 4: Per-query vs. scalar- $\hat{\sigma}$ ablation, IV-surface branch, architecture A, seed 42. The scalar- $\hat{\sigma}$ row is a restricted control, not a candidate model. The per-query rows repeat the canonical Table 2 values; the separate ablation evaluator returns 0.221896, 0.977899, and 0.294127 for the three cells where the two scripts differ. The scalar-control rows come from the independent CPU ablation evaluator in float32 at batch size 4096 over the full test split. As a cross-check, that evaluator reproduces the per-query results to the fifth decimal, with the three displayed rounding differences noted above.

Option	$\hat{\sigma}$ conditioning	$R^2(\text{price})$	$R^2(\log)$	RMSE(log)
Call	per-query	0.999977	0.991044	0.221894
Call	scalar (market state only)	0.999825	0.800103	1.048319
Put	per-query	0.992553	0.977898	0.294136
Put	scalar (market state only)	-3.090190	-0.316871	2.270415

On calls, architectures A and B are effectively tied (0.990996 vs. 0.991075). The mean difference slightly exceeds architecture A’s across-seed standard deviation but lies well within architecture B’s: 0.000080 against 0.000073 for A and 0.000716 for B. On puts, architecture B’s mean is somewhat higher than architecture A’s (0.981738 vs. 0.979099), but this reading rests on only three seeds and is not, by itself, a strong superiority claim. We therefore do not rank the two architectures: the observed differences are small relative to the measured seed-to-seed variation, and three seeds are insufficient to rank them. We ran no inferential test, and none of the comparisons in this subsection should be read as a significance test; the $\pm$ figures in Table 3 are descriptive across-seed standard deviations over three runs, nothing more. The SPX-history and VIX-history branches have no replication at all—one seed each—so no dispersion estimate exists for the other eight cells of Table 2. For those branches we cannot say whether an observed A-versus-B difference is larger or smaller than run-to-run noise, and we therefore draw no architecture comparison there at all.

6.3 Ablation: per-query volatility conditioning

Table 4 isolates the effect of per-query $\hat{\sigma}$ conditioning by comparing the full IV-surface architecture-A model against the scalar- $\hat{\sigma}$ control of Section 5.3.

Removing per-contract conditioning drops call $R^2(\log)$ from 0.991 to 0.800 and put $R^2(\log)$ from 0.978 to about -0.32 ; the scalar- $\hat{\sigma}$ log-space figures are quoted to two decimals because that is the precision they reproduce to across evaluation paths (Appendix B). On puts, the price-space effect is larger still: $R^2(\text{price})$ falls from 0.993 to -3.090 , worse than always predicting the mean price. Both rows are single-seed (seed 42, Section 5.3), so the exact figure -3.090 describes this one run and should not be read as a replicated estimate. What is seed-independent here is the representational restriction: a market-state-only $\hat{\sigma}$ is constant across all contracts sharing a market state and cannot express a smile or skew, a property of the architecture alone. The direction and magnitude of the observed test metrics, including the fact that they turn the put figures negative at all, come from a single unreplicated seed and are not guaranteed by that restriction. Figure 2 shows the visual counterpart of this result on the volatility surface itself: at three fixed test dates and three maturity buckets per option type, the per-query $\hat{\sigma}$ curve tracks the observed market-IV smile/skew across moneyness, while the scalar- $\hat{\sigma}$ curve is, by construction, flat within each panel. Both models here take the IV-surface branch, so the per-query curve demonstrates that the architecture can reconstruct the day’s smile from that day’s standardized surface plus the contract query (Section 4.1); it is not evidence that the smile is recovered from information available strictly

Table 5: Calls: baselines are close to saturated on $R^2(\text{price})$; all five numbers lie within a 0.0014 range, so this metric alone does not discriminate pricers on this dataset. The log-space column separates them sharply. Both constant- $\hat{\sigma}$ rows are train-subsample-fitted (Section 5.2), one by minimizing MSE on $\log(\frac{V}{K})$ and one by minimizing MSE on $\frac{V}{K}$.

Baseline / model	$R^2(\text{price})$	$R^2(\log)$
Zero-volatility limit (discounted forward intrinsic)	0.998586	-16.806016
Constant $\hat{\sigma}$, log-MSE fit ($\sigma=0.1933$)	0.999706	0.620788
Constant $\hat{\sigma}$, price-MSE fit ($\sigma=0.1761$)	0.999800	0.640351
IV-surface spline interpolation	0.999982	0.764181
Best neural model (IV surface, A)	0.999977	0.991044

Table 6: Puts: price-space and log-space metrics disagree on which pricer wins. Both constant- $\hat{\sigma}$ rows are train-subsample-fitted (Section 5.2), one by minimizing MSE on $\log(\frac{V}{K})$ and one by minimizing MSE on $\frac{V}{K}$; the two objectives select different constants and rank in opposite directions on the two metrics.

Baseline / model	$R^2(\text{price})$	$R^2(\log)$
Zero-volatility limit (discounted forward intrinsic)	0.535219	-34.203142
Constant $\hat{\sigma}$, log-MSE fit ($\sigma=0.4758$)	-2.358206	-0.388064
Constant $\hat{\sigma}$, price-MSE fit ($\sigma=0.2063$)	0.925761	-3.358130
IV-surface spline interpolation	0.996594	-2.518559
Best neural model (IV surface, B)	0.993780	0.981298

before the quote.

6.4 Error stratification by moneyness and maturity

Figure 3 reports $\text{RMSE}(\log \frac{V}{K})$ over a maturity $\times$ log-moneyness grid on the full test split for the IV-surface, architecture-A models. Both option types show the same qualitative pattern: error is highest in the deep out-of-the-money wings at short maturity and lowest near the money at medium-to-long maturity. This residual near-expiry effect reflects a genuine, continuous ill-conditioning of the inverse map $\log(\text{price}) \rightarrow \sigma$ as $T \rightarrow 0$, since the sensitivity $\frac{\partial \log(\text{price})}{\partial \sigma}$ shrinks near expiry and away from the money, not a degenerate or unfittable region of the input: this dataset contains no contracts with $T \leq 1$ day by construction, so the shortest surviving maturity band is simply where this continuous effect is most visible, and on the price scale (rather than the log scale) the same cells are small in absolute terms. The maturity-stratified rows of Figure 3 are how we report this near-expiry pattern; we do not present a separate maturity-resolved tail plot.

6.5 Comparison against baselines

Table 5 and Table 6 compare the best model against the three baselines of Section 5.2.

On calls, all five numbers in Table 5 sit within a 0.0014 range: the zero-volatility limit, with no fitted parameters, already reaches $R^2(\text{price}) = 0.9986$, so call price- R^2 is a weak discriminator on its own and the log-space metric is the more informative one for calls. That column separates the same pricers cleanly: -16.81 for the zero-volatility limit, 0.621 and 0.640 for the log-fitted and price-fitted constant $\hat{\sigma}$, 0.764 for spline interpolation of the observed surface, and 0.991 for the model. Unlike puts, calls show no material metric disagreement: the model leads every baseline


on log-space and sits within 6×10^{-6} of the best baseline on price-space (a gap of 5.5×10^{-6}). On puts, the two metrics disagree on which pricer wins: IV-surface spline interpolation edges out the model on price-space (0.996594 vs. 0.993780), while the model leads on log-space by 3.50 in R^2 (0.981298 vs. -2.518559). Whichever pricer a reader considers “better” for puts depends on which metric they weight; we report both rather than only the more flattering one.

The spline’s negative put $\log\text{-}R^2$ is a grid-boundary artefact rather than a general failure of surface interpolation. A row-level decomposition of the same baseline (which reproduces the aggregate R^2 of Table 6 to 10^{-6}) finds that 578,610 of the 1,380,379 put test rows (41.92%) converge to $|\delta|$ below the surface’s 10-delta boundary and are therefore all priced off that one clamped column. Those rows carry 99.94% of the spline’s log-space sum of squared errors and all 272,433 squared-log errors above 10, while the rows clamped on neither axis reach $R^2(\log) = 0.9939$ (0.9843 on the looser reading that excludes only the delta clamp and keeps tenor-clamped rows). The mechanism is a scale mismatch between the two metrics: clamping a deep out-of-the-money put to the 10-delta implied volatility is a large multiplicative error, but it acts on a small normalized price, the mean true $\frac{V}{K}$ on those rows being 0.00203 against 0.0202 in-grid and 0.0886 deep in the money. They still carry 72.94% of the price-space sum of squared errors, a substantial share and not a negligible one, but $R^2(\text{price})$ is scored against a variance denominator that the in-the-money rows set, and that denominator is large enough to leave it at 0.996594. Calls show the same structure more weakly, with 185,863 rows (22.68%) below 10 delta carrying 99.40% of call log-space SSE. This is also where the learned model’s log-space advantage over interpolation is concentrated, since the standardized surface has no support beyond 10 delta; the model’s own error heatmap (Figure 3) places its worst cells in the same wings, so the advantage there is relative rather than absolute.

The put constant- $\hat{\sigma}$ baseline depends strongly on the objective it was fitted under, so what ranks it is the choice of fitting space and not the constancy of $\hat{\sigma}$ alone. Fitted by minimizing MSE on $\log(\frac{V}{K})$ it selects $\sigma = 0.4758$ and scores below the zero-volatility limit on put price-space (-2.358206 vs. 0.535219); fitted by minimizing MSE on $\frac{V}{K}$ it selects $\sigma = 0.2063$ and reaches $R^2(\text{price}) = 0.925761$, but falls to $R^2(\log) = -3.358130$. The log-space objective is dominated by deep out-of-the-money puts, whose implied volatilities are the highest in the cross-section, so the constant it selects is far too high for near-the-money puts and overprices them badly on the price scale, where those contracts carry almost all of the variance. The price-fitted constant reverses the picture, pricing the near-the-money contracts acceptably at the cost of the wings that dominate the log metric. The claim is therefore narrow: it is the log-fitted constant that scores below the zero-volatility limit on put price-space, not any single global volatility.


Figure 4 shows predicted-vs-market price scatter on a 50,000-row uniform sample of the test split. Calls reach $R^2(\text{price}) = 0.999954$, Pearson correlation (log-log) $r = 0.995619$, median relative error 3.61%, and 76.62% of rows within $\pm 10\%$. Puts reach $R^2(\text{price}) = 0.993179$, $r = 0.991166$, median relative error 8.74%, and 54.46% of rows within $\pm 10\%$.

F2 — smile recovery, calls (v5 test split): per-query $\hat{\sigma}$ tracks the market smile, scalar- σ cannot




(a) Calls


F2 — smile recovery, puts (v5 test split): per-query $\hat{\sigma}$ tracks the market smile, scalar- σ cannot



(b) Puts



(a) Calls



(b) Puts

Figure 3: $\text{RMSE}(\log \frac{V}{K})$ over a 5×9 maturity $\times$ log-moneyness grid, full test split, IV-surface architecture-A models. Error is highest in the deep out-of-the-money wings at short maturity and lowest near-the-money at longer maturity; the two grids mirror each other because a call's out-of-the-money side is the opposite sign of log-moneyness from a put's.

Table 7: Decoder derivative-consistency check: autograd vs. closed-form BS partial greeks at the predicted $\hat{\sigma}$, held fixed. $n=2,000$ sampled test contracts per option type.

Option	Delta max abs. diff.	Delta median abs. diff.	Vega max abs. diff.	Vega median abs. diff.
Call	1.84×10^{-14}	1.40×10^{-16}	1.38×10^{-15}	5.55×10^{-17}
Put	1.43×10^{-14}	8.37×10^{-17}	2.12×10^{-15}	5.55×10^{-17}




Figure 4: Predicted vs. market normalized price ($\frac{V}{K}$), log-log axes, colored by maturity bucket, 50,000-row uniform sample of the test split, IV-surface architecture-A models.

6.6 Black–Scholes decoder derivative consistency

Table 7 reports the agreement between autograd-computed and closed-form delta and vega at the predicted $\hat{\sigma}$, over 2,000 sampled test contracts per option type.

Agreement is at float64 machine precision, as expected for two mathematically identical expressions evaluated by two different computational paths. The largest absolute discrepancy over the sampled contracts is 1.84×10^{-14} for call delta and 1.43×10^{-14} for put delta, and 1.38×10^{-15} for call vega and 2.12×10^{-15} for put vega, with medians at or below 1.40×10^{-16} ; we quote the two greeks separately rather than as a single range, since their maxima differ by an order of magnitude. This confirms the decoder is implemented and differentiated correctly at the predicted $\hat{\sigma}$; it is not an independent check of the learned $\hat{\sigma}$ itself, and it does not measure the total (sticky-strike) sensitivity of the learned system, which would include a $\frac{\partial \hat{\sigma}}{\partial S} \cdot \text{vega}$ term this frozen- $\hat{\sigma}$ check does not exercise.

7 Limitations

The results above are bounded by several scope limits that we state explicitly rather than leave implicit.

Pointwise, not surface-wide, Black–Scholes consistency. At a fixed predicted $\hat{\sigma}$, the decoder guarantees the single-contract static no-arbitrage bounds at each queried contract, but not calendar-spread or butterfly consistency across the predicted surface, and not that the total learned map, in which $\hat{\sigma}$ varies with the query, satisfies the BS PDE. A violation-rate diagnostic on the predicted (K, T) surface would be needed to make a surface-wide claim, and we do not run one here.

Call price- R^2 is a weak metric on this dataset. The zero-volatility limit (discounted forward intrinsic), which fits no parameter, already reaches $R^2(\text{price}) = 0.9986$ for calls (Table 5), so most of the call price- R^2 figures in this paper should not be read as strong evidence on their own; the put comparison and the log-space metric carry the informative signal.

The put interpolation comparison is shaped by grid clamping. The spline baseline’s negative put log- R^2 comes almost entirely from rows clamped to the standardized surface’s 10-delta boundary (Section 6.5), so it measures a limit of that grid’s support rather than an absence of information in the observed surface. Reading the model’s log-space lead over interpolation as a statement about surface-based pricing in general would require a comparator that extrapolates beyond the grid, such as an SVI or SSVI fit [9, 10], which we do not evaluate here.

Architecture ranking (A vs. B) is underdetermined. Three-seed replication exists only for the four IV-surface cells, and even there the put result, B somewhat ahead of A, rests on three seeds, not enough for a superiority claim; we report descriptive across-seed dispersion and run no inferential test anywhere in this paper. The SPX-history and VIX-history branches are single-seed throughout, with genuinely unmeasured dispersion, so we make no architecture claim for those eight cells. The scalar- $\hat{\sigma}$ ablation is likewise single-seed (Section 5.3), so its magnitudes describe one run.

The surface branch is conditioned on same-day information. Its input is OptionMetrics’ own standardized surface for the trade date being priced, built from that day’s quotes (Section 4.1). The date-based split removes cross-date leakage but not this within-date dependence, so the surface-branch results establish same-day cross-sectional reconstruction conditional on an observed surface, not forecasting or independent price discovery; any deployment reading of those numbers would require the surface to be available at pricing time. The historical branches are not strictly-prior-information inputs either: their 21-trading-day windows end on the target trade date and are normalized by that day’s close, and VIX is itself computed from contemporaneous SPX option quotations. Their same-day option-market content is far smaller and less direct than the surface branch’s, but a design that required strictly pre-trade-date information would have to lag these windows by a day, which we did not do.

The branch comparison compares configurations, not inputs, and the VIX representation is level-free. The first, grid-dependent layer of the two-layer flattening projection holds 766,080 (surface), 430,208 (SPX history), and 344,192 (VIX history) parameters, a spread of roughly 2.2 $\times$; grid size, spectral modes, input timing, and normalization also differ across the three branches. Input content is therefore confounded with all of these in Section 6.1, and the result should be read as a ranking of three whole input–encoder configurations rather than an isolation of input content. Separately, the VIX-history tensor is divided by the window’s final VIX close, which discards the

absolute VIX level and retains only the path shape, so the VIX result speaks only to that close-normalized representation and not to VIX history as an information source. A capacity-matched rerun with a level-preserving VIX encoding would be needed to attribute the branch gap to input content alone.

Near-expiry ill-conditioning is real and not removed by the maturity filter. The maturity filter removes the degenerate exact- $T=0$ case but not the continuous ill-conditioning of $\frac{\partial \log(\text{price})}{\partial \hat{\sigma}}$ as $T \rightarrow 0$; the shortest surviving maturity band still shows elevated log-space error (Section 6.4), though it is small in absolute price terms.

Single underlying, European exercise only. All results are on SPX index options. European exercise removes the early-exercise complication, and BS is used as the analytical decoder, with the dividend yield handled through the decoder’s q term. We make no claim about American-style contracts or other underlyings, and the remaining BS idealizations, constant volatility over the contract’s life and frictionless continuous trading, are not satisfied by the market; this is precisely why $\hat{\sigma}$ must be predicted per contract in the first place.

Mid-price only. We price the bid-ask midpoint; transaction costs and the bid-ask spread itself are out of scope.

No direct-price-regressor control. We do not compare against a direct-price regressor of comparable size that skips the BS decoder entirely, so this paper makes no claim about how the BS-decoder design compares to that alternative.

8 Conclusion

Predicting a per-contract implied volatility and pricing it through a differentiable Black–Scholes decoder removes the need for a learned pricing surrogate. Conditional on the predicted $\hat{\sigma}$ being held fixed, this yields two structural properties by construction: analytic partial greeks at that $\hat{\sigma}$, with autograd through the decoder and closed-form BS agreeing to a maximum absolute discrepancy of approximately 2×10^{-14} for delta and 2×10^{-15} for vega (Section 6.6), and pointwise Black–Scholes-consistent prices at every queried contract (Section 3.4). Neither extends to the total learned map, whose $\hat{\sigma}$ varies with the query, nor to arbitrage-freedom across the full predicted surface.

The per-query-vs-scalar ablation (Section 6.3) shows that per-query conditioning matters: removing it collapses put $R^2(\log)$ from 0.978 to about -0.32 and put $R^2(\text{price})$ to -3.09 , worse than predicting the mean. Comparing three candidate market-state input–encoder configurations, the one reading the contemporaneous implied-volatility surface consistently prices better than either historical configuration, on both option types and both head architectures (Section 6.1). The surface-branch models were stable across three seeds, and the observed seed-42 gap between the surface configuration and the historical ones is much larger than the surface branch’s own seed-to-seed dispersion; full robustness of the branch-versus-branch comparison cannot be established without also replicating the SPX-history and VIX-history runs, which remain single-seed. That comparison is not, however, like-for-like: the surface branch reads OptionMetrics’ own same-day standardized surface, so it demonstrates out-of-time (date-split), same-day cross-sectional reconstruction conditional on an observed surface rather than forecasting, and the three encoders differ in grid size, normalization, and parameter count while the VIX input is normalized in a way that

removes the absolute VIX level (Section 7). The separate question of which per-query head architecture is better remains unresolved: the observed differences are small relative to the seed-to-seed variation we measure, and three seeds are insufficient to rank them.

Taken together, on this dataset the approach reaches $R^2(\text{price}) \approx 0.9999\text{--}0.99998$ and $R^2(\log) \approx 0.990\text{--}0.991$ on the IV-surface input for calls, though calls’ price- R^2 is close to the zero-volatility limit and should not be read in isolation. The more informative put comparison shows the model 0.003–0.004 below IV-surface interpolation on price-space (0.993–0.994 vs. 0.997) but ahead of it by roughly 3.5 in $R^2(\log)$ (0.978–0.981 vs. -2.52), a gap that the grid-boundary decomposition of Section 6.5 attributes to clamping at the surface’s 10-delta edge.

References

- [1] Damien Ackerer, Natasa Tagasovska, and Thibault Vatter. Deep smoothing of the implied volatility surface. In *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*, pages 11552–11563, 2020.
- [2] Fischer Black and Myron Scholes. The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3):637–654, 1973.
- [3] Matthias Büchner and Bryan T. Kelly. A factor model for option returns. *Journal of Financial Economics*, 143(3):1140–1161, 2022. doi: 10.1016/j.jfineco.2021.12.007.
- [4] Cboe Global Markets. VIX index: Frequently asked questions. https://www.cboe.com/tradable_products/vix/faqs, 2026. Accessed 2026.
- [5] Samuel N. Cohen, Christoph Reisinger, and Sheng Wang. Detecting and repairing arbitrage in traded option prices. *Applied Mathematical Finance*, 27(5):345–373, 2020. doi: 10.1080/1350486X.2020.1846573.
- [6] Robert Culkin and Sanjiv R. Das. Machine learning in finance: The case of deep learning for option pricing. *Journal of Investment Management*, 15(4):92–100, 2017.
- [7] Lijie Ding, Egang Lu, and Kin Cheung. Deep learning option pricing with market implied volatility surfaces. arXiv preprint, 2025.
- [8] Bruno Dupire. Pricing with a smile. *Risk*, 7(1):18–20, 1994.
- [9] Jim Gatheral. A parsimonious arbitrage-free implied volatility parameterization with application to the valuation of volatility derivatives. Presentation, Global Derivatives & Risk Management, Madrid, 2004.
- [10] Jim Gatheral and Antoine Jacquier. Arbitrage-free SVI volatility surfaces. *Quantitative Finance*, 14(1):59–71, 2014. doi: 10.1080/14697688.2013.819986.
- [11] Lukas Gonon, Antoine Jacquier, and Ruben Wiedemann. Operator deep smoothing for implied volatility. In *International Conference on Learning Representations (ICLR)*, 2025.
- [12] Blanka Horvath, Aitor Muguruza, and Mehdi Tomas. Deep learning volatility: A deep neural network perspective on pricing and calibration in (rough) volatility models. *Quantitative Finance*, 21(1):11–27, 2021. doi: 10.1080/14697688.2020.1817974.

- [13] James M. Hutchinson, Andrew W. Lo, and Tomaso Poggio. A nonparametric approach to pricing and hedging derivative securities via learning networks. *The Journal of Finance*, 49(3):851–889, 1994. doi: 10.1111/j.1540-6261.1994.tb00081.x.
- [14] Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Fourier neural operator for parametric partial differential equations. In *International Conference on Learning Representations (ICLR)*, 2021.
- [15] Lu Lu, Pengzhan Jin, Guofei Pang, Zhongqiang Zhang, and George Em Karniadakis. Learning nonlinear operators via DeepONet based on the universal approximation theorem of operators. *Nature Machine Intelligence*, 3:218–229, 2021. doi: 10.1038/s42256-021-00302-5.
- [16] Robert C. Merton. Theory of rational option pricing. *The Bell Journal of Economics and Management Science*, 4(1):141–183, 1973.
- [17] Miquel Noguer i Alonso. Derivative-informed operator learning for finance: On-the-fly greeks, surfaces, hedging, and control. arXiv preprint, 2026.
- [18] OptionMetrics. IvyDB US: Standardized options and volatility surface data. Product documentation, <https://optionmetrics.com/data-products/>, 2026. Accessed 2026.
- [19] Maziar Raissi, Paris Perdikaris, and George Em Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019. doi: 10.1016/j.jcp.2018.10.045.
- [20] Johannes Ruf and Weiguan Wang. Neural networks for option pricing and hedging: A literature review. *Journal of Computational Finance*, 24(1):1–46, 2020. doi: 10.21314/JCF.2020.390.

A Hyperparameters

Identical across all twelve core runs: seed 42 (plus seeds 43, 44 for the four IV-surface replication runs); Adam phase, 100 epochs, learning rate 10^{-3} with reduce-on-plateau scheduling; fine-tuning phase, 50 epochs, learning rate 10^{-5} , minimum 10 fine-tuning epochs; batch size 256, gradient-norm clipping at 1.0; early-stopping patience 5 in each phase. The FNO encoder uses 4 spectral convolution blocks with 1×1 -convolution skip connections and SiLU activations, width 32, flattened and projected to a 128-dimensional latent/basis vector; head/trunk hidden width is 128. The number of retained spectral modes follows the grid: 4×8 (modes1 $\times$ modes2) for the surface encoder on its 11×17 grid, and 6×2 for both the SPX-history (21×5) and VIX-history (21×4) encoders. Architecture B (DeepONet) additionally uses a small positive $\hat{\sigma}$ floor and a learned scalar bias initialized so that $\hat{\sigma} \approx 0.20$ at initialization.

B Numerical stability

The default training and evaluation path computes $\log(\text{clamp}(\frac{V}{K}, 10^{-8}))$ for numerical stability. A numerically exact alternative, using a stable log-normal-CDF formulation, is available for evaluation only: its gradient is exact but explodes near the money, which causes training to diverge if used during optimization, so it is used only as a read-only diagnostic, never during training.

A controlled re-evaluation of the same scalar- $\hat{\sigma}$ checkpoints on CPU (batch 512 and 4096) and GPU (batch 512, the training-time path) reproduces $R^2(\text{price})$ to within 4×10^{-7} but $R^2(\log)$ only to about 3×10^{-3} across device paths, and decoding the same float32 $\hat{\sigma}$ in float64 moves $R^2(\log)$ further, to 0.798 for calls and -0.328 for puts, because cancellation in the float32 Black-Scholes decoder, before the 10^{-8} price clamp and logarithm, moves near-zero predictions on and off the clamp floor across device paths and precisions; the per-query models, which never reach the floor, differ by at most 1.5×10^{-6} in $R^2(\log)$ across the evaluated paths and precisions (`results/ablation_evaluator_consistency_v5.json`). The scalar- $\hat{\sigma}$ log-space metrics are therefore stable only to about 10^{-2} in $R^2(\log)$ for puts and 2×10^{-3} for calls, which is why the prose reports them to two decimals; Table 4 keeps the canonical float32 CPU evaluator’s digits. No implementation error was found.

C Reproducibility and environment

Table 8 records the environment captured on 2026-09-07 by `analysis/snapshot_environment.py` (`results/environment_snapshot_2026-09-07.txt`) and used for the publication-audit diagnostics described above. The August 2026 training environment was not snapshotted, and this table should not be interpreted as reconstructing either the training environment or every earlier evaluation run.

Table 8: Environment captured 2026-09-07 and used for the publication-audit diagnostics. It is neither the training environment nor necessarily that of every earlier evaluation run.

Component	Version
Python	3.11.15
PyTorch	2.10.0 (built against CUDA 13.0)
NumPy	2.3.5
h5py	3.16.0
matplotlib	3.10.8
polars	1.38.1
GPU	NVIDIA GeForce RTX 3060
GPU driver	595.95

Each run’s exact configuration is saved alongside its checkpoints at training time; that record covers the seed, the hyperparameters, and the dataset version and checksum, but not package versions.